

6th International Conference on
I-SMAC (IoT in Social, Mobile, Analytics and Cloud)
I-SMAC 2022

**A REVIEW ON STATISTICAL POWER MODELLING
FOR A GRAPHICS PROCESSING UNIT (GPU)**

Presented By

Yojan Chitkara



Table of Contents

- INTRODUCTION
- GPU BASICS
- SILICON POWER
- RESULTS
- REFERENCES



INTRODUCTION



Scaling Architectures

- GPU computing is defining a new Supercharged Law, specifically due to the huge parallelism in its computing architecture.




Figure 1.1: GPU Parallelism Compute Perf [7]



Shifting Trends towards GPU

- The Graphics Processing Unit (GPU) or the Virtual Processing Unit (VPU) as shown in Figure 1.2 is a specialized electronic circuit designed to rapidly manipulate or alter memory to accelerate heavy workloads in shorter intervals of time as compared to a CPU.




Figure 1.2: Parallelism between CPU and GPU architectures. [2]

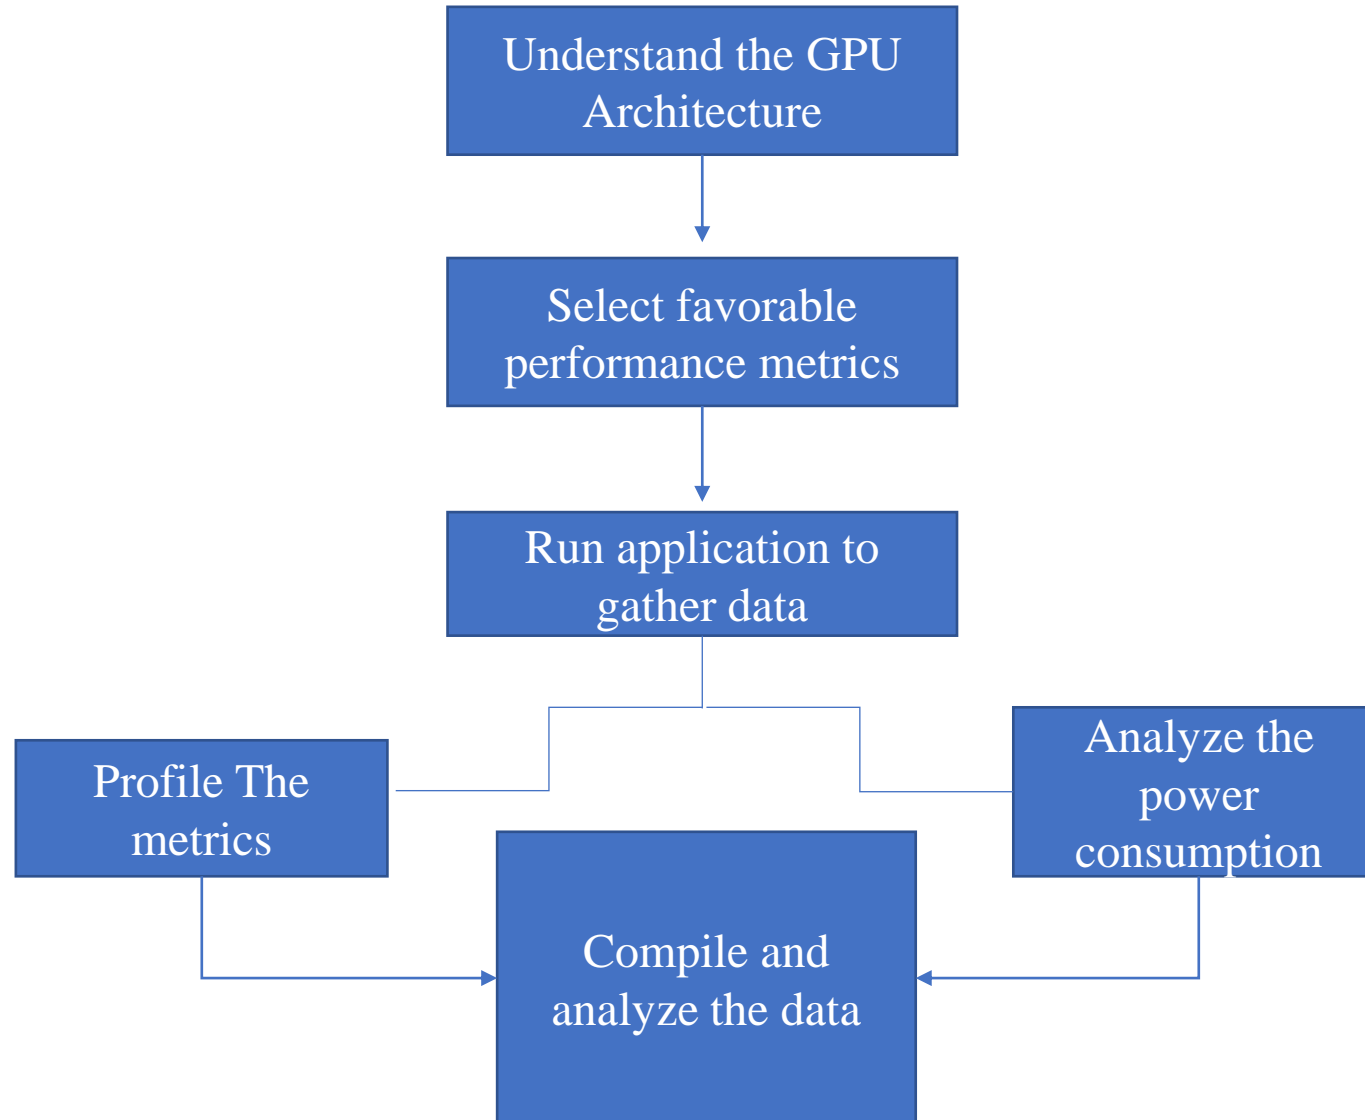


Motivation

- Designing for low power as architectures scaled down in technology became essential as number of transistors on a die kept on increasing to accommodate more cores.
- It is hence becoming increasingly difficult for speed to keep par with increasing number of transistors.



Working Methodology



GPU BASICS



Parallelism in CPU's

- The five essential steps required for an instruction to execute in a CPU RISC Architecture are:
 - Instruction Fetch (IF)
 - Instruction Decode (ID)
 - Instruction Execute (Ex)
 - Memory Access (Mem)
 - Register Write-Back (WB)



Pipelining

- Pipelining is a technique of executing multiple instructions in parallel per clock. Instruction Level Parallelism can fix the inefficiencies brought about by sequential execution like latency, resource utilization.

Instruction No	Pipeline Stage						
1	IF	ID	Ex	Mem	WB		
2		IF	ID	Ex	Mem	WB	
3			IF	ID	Ex	Mem	WB
4				IF	ID	Ex	Mem
5					IF	ID	Ex
Clock Cycle	1	2	3	4	5	6	7



GPU Design

- A GPU comprises of a large number of cores (significantly greater than a CPU) each of which run at a clock speed significantly slower than a core of a CPU.


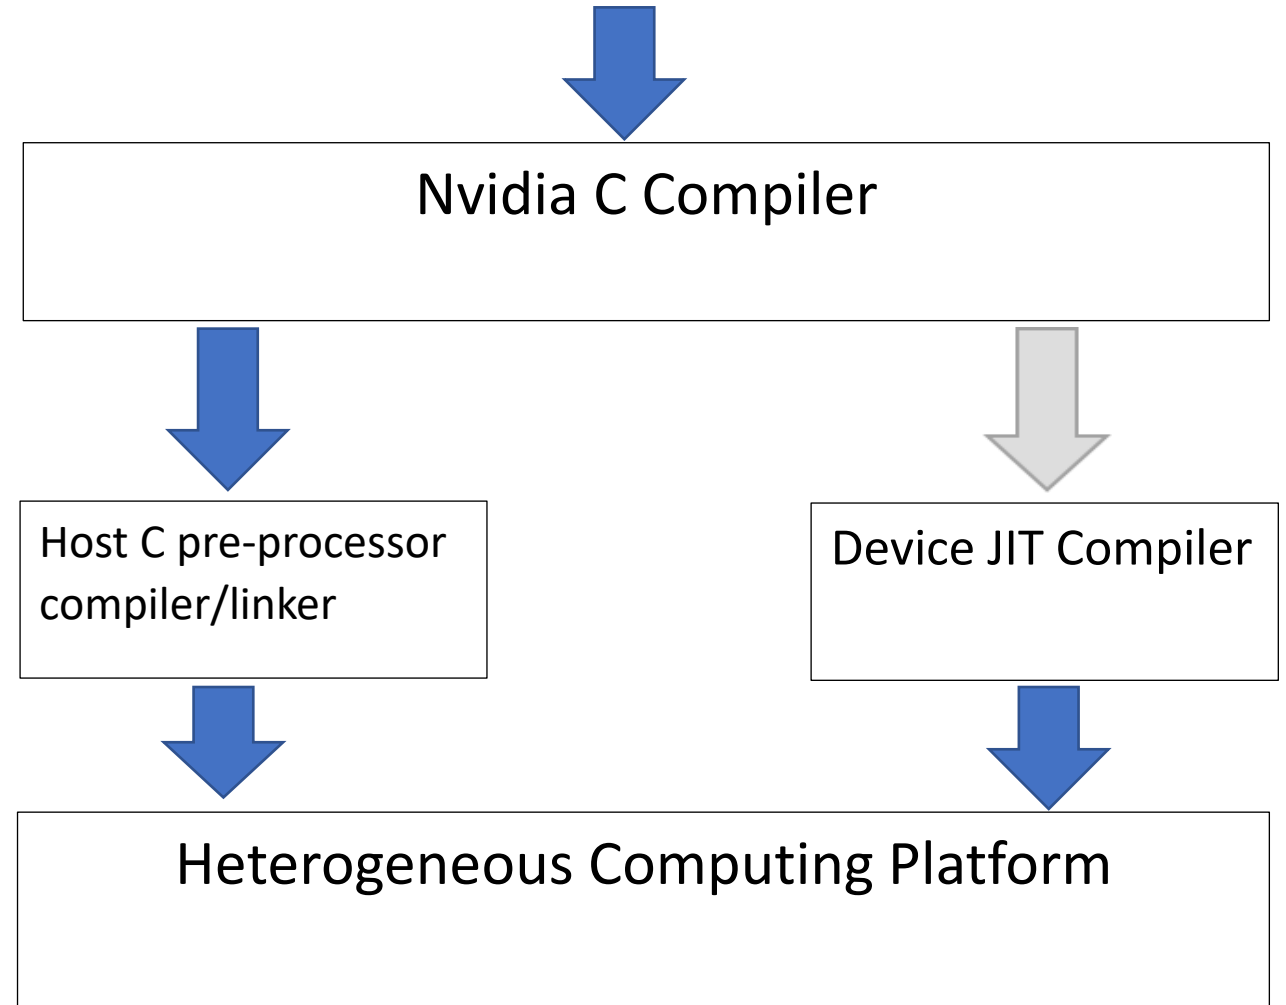


Figure 1.3 : GPU Architecture



GPU Execution

- A typical CUDA program has code intended for both the CPU and the GPU
- The NVCC or Nvidia C Compiler splits the Data intensive code from the program.
- The Host machine processes the less intensive code.
- The Device processes the more intensive code.
- Together it forms a heterogeneous computing platform.



CUDA Memories

- Following are the supported memory types on a CUDA supported GPU
 - R/W per-thread registers
 - R/W per-thread local memory
 - R/W per-block shared memory
 - R/W per-grid global memory
 - Read only per-grid constant memory.




Figure 1.4: Thread Level Hierarchy



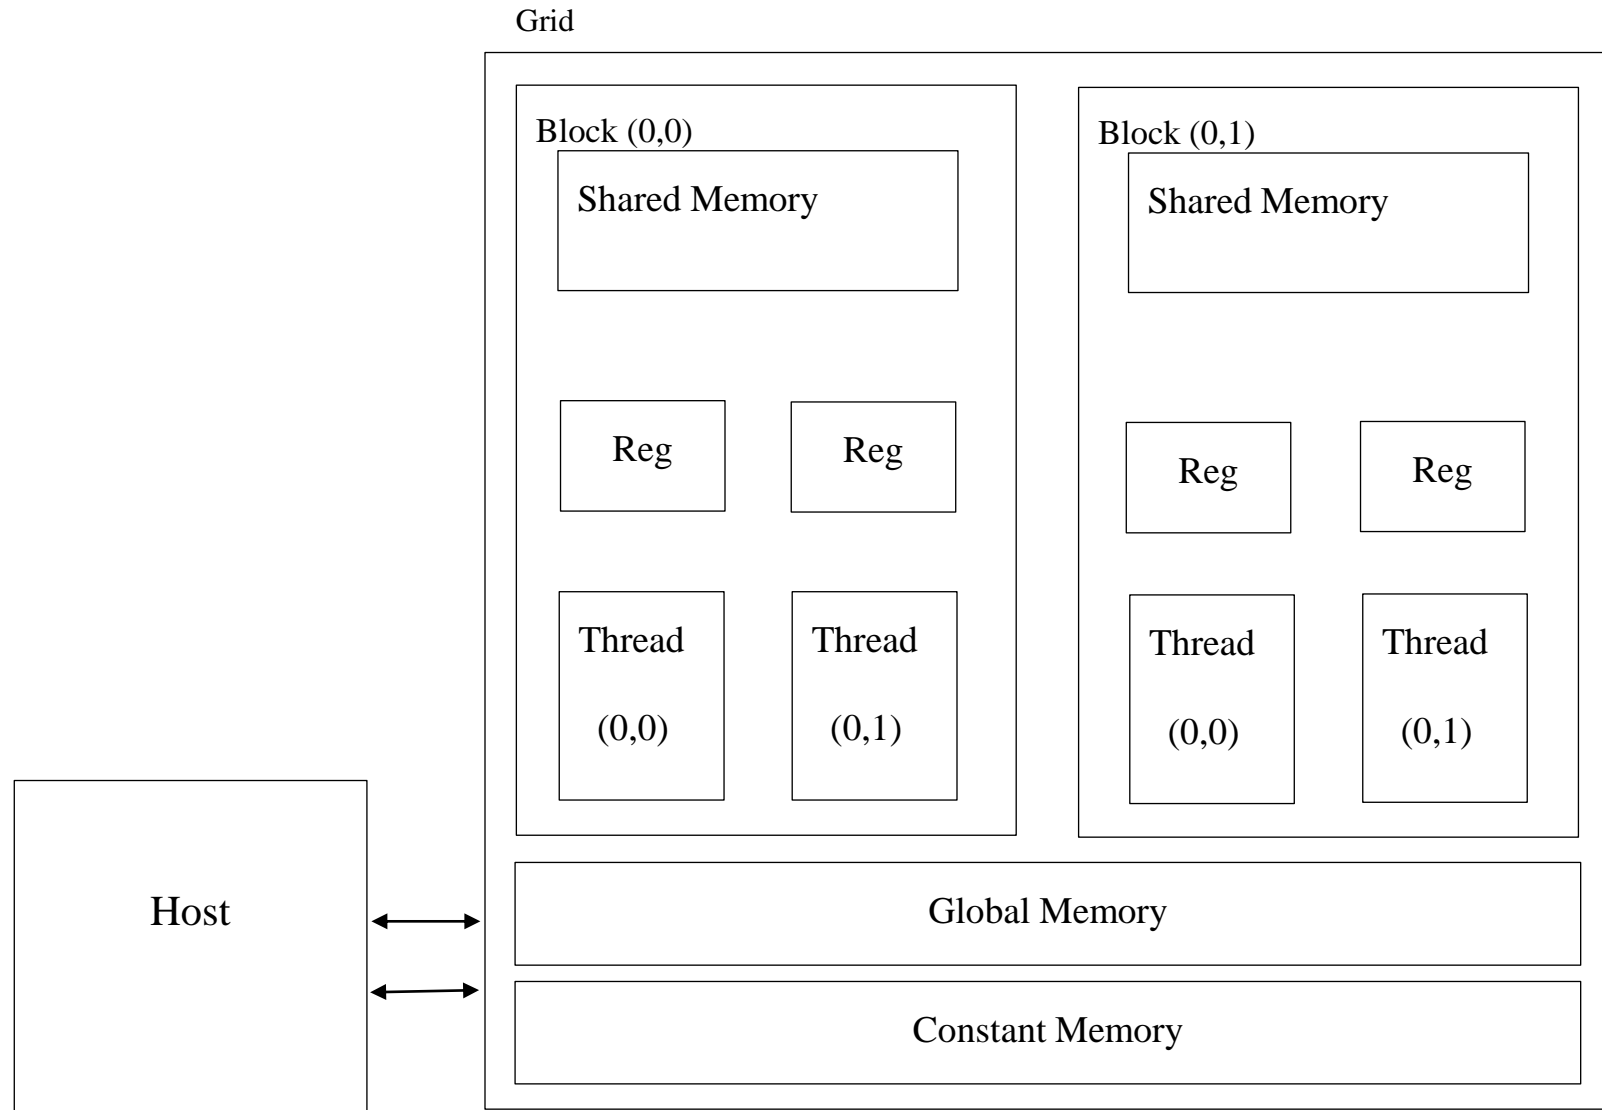


Figure 1.5: Memory Assignment



SILICON POWER



Dynamic Power

- Consists of the power consumed during logic transitions and can be sub-divided into two types, Switching and Internal Power.
 - Switching Power: As the name suggests, this power manifests itself when there are transitions between logics.

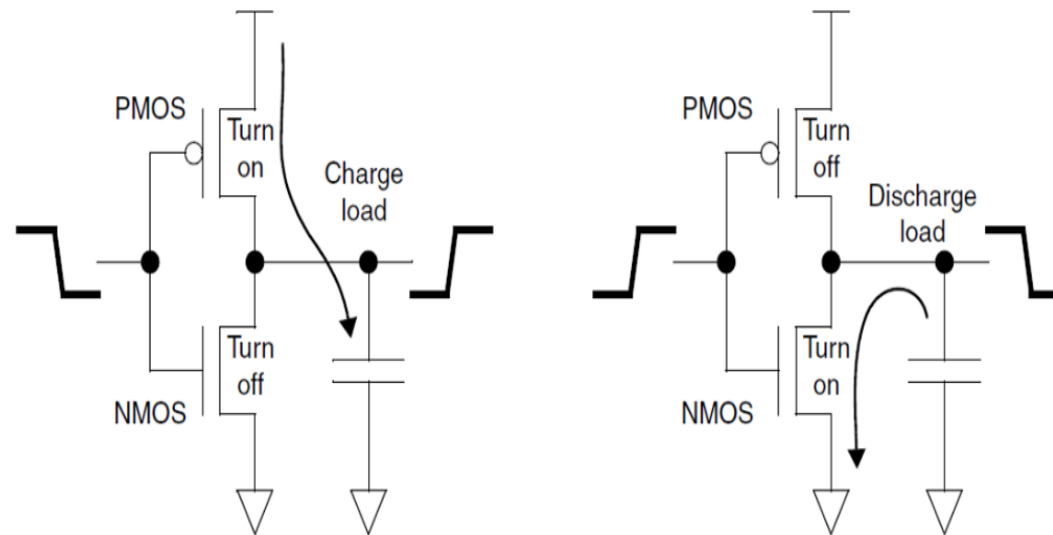


Figure 1.6: Logic transitions in an inverter



- Internal Power: When the input signal is at an intermediate voltage level, a relatively large amount of current flows through the transistors for a brief period of time as shown in Figure 1.7 contributing to internal power.

$$P_{dyn} = (C_{eff} \cdot V_{dd}^2 \cdot f_{clk}) + (t_{sc} \cdot V_{dd} \cdot I_{peak} \cdot f_{clk}) \quad (1.1)$$




Figure 1.7: Short Circuit Current



Static (Leakage) Power

- Static power is the transistor leakage current whenever power is applied to the device. The main causes of leakage power as displayed in Figure 1.8 are
 - Reverse bias p-n junction diode leakage: Caused by minority carrier drift
 - Sub threshold leakage: Current which flows from drain to source when transistor operates in weak inversion region. Sub threshold leakage is the biggest contributor of leakage power.
 - Gate leakage: Current which flows directly from the gate to the substrate either through hot carrier injection or oxide tunneling.



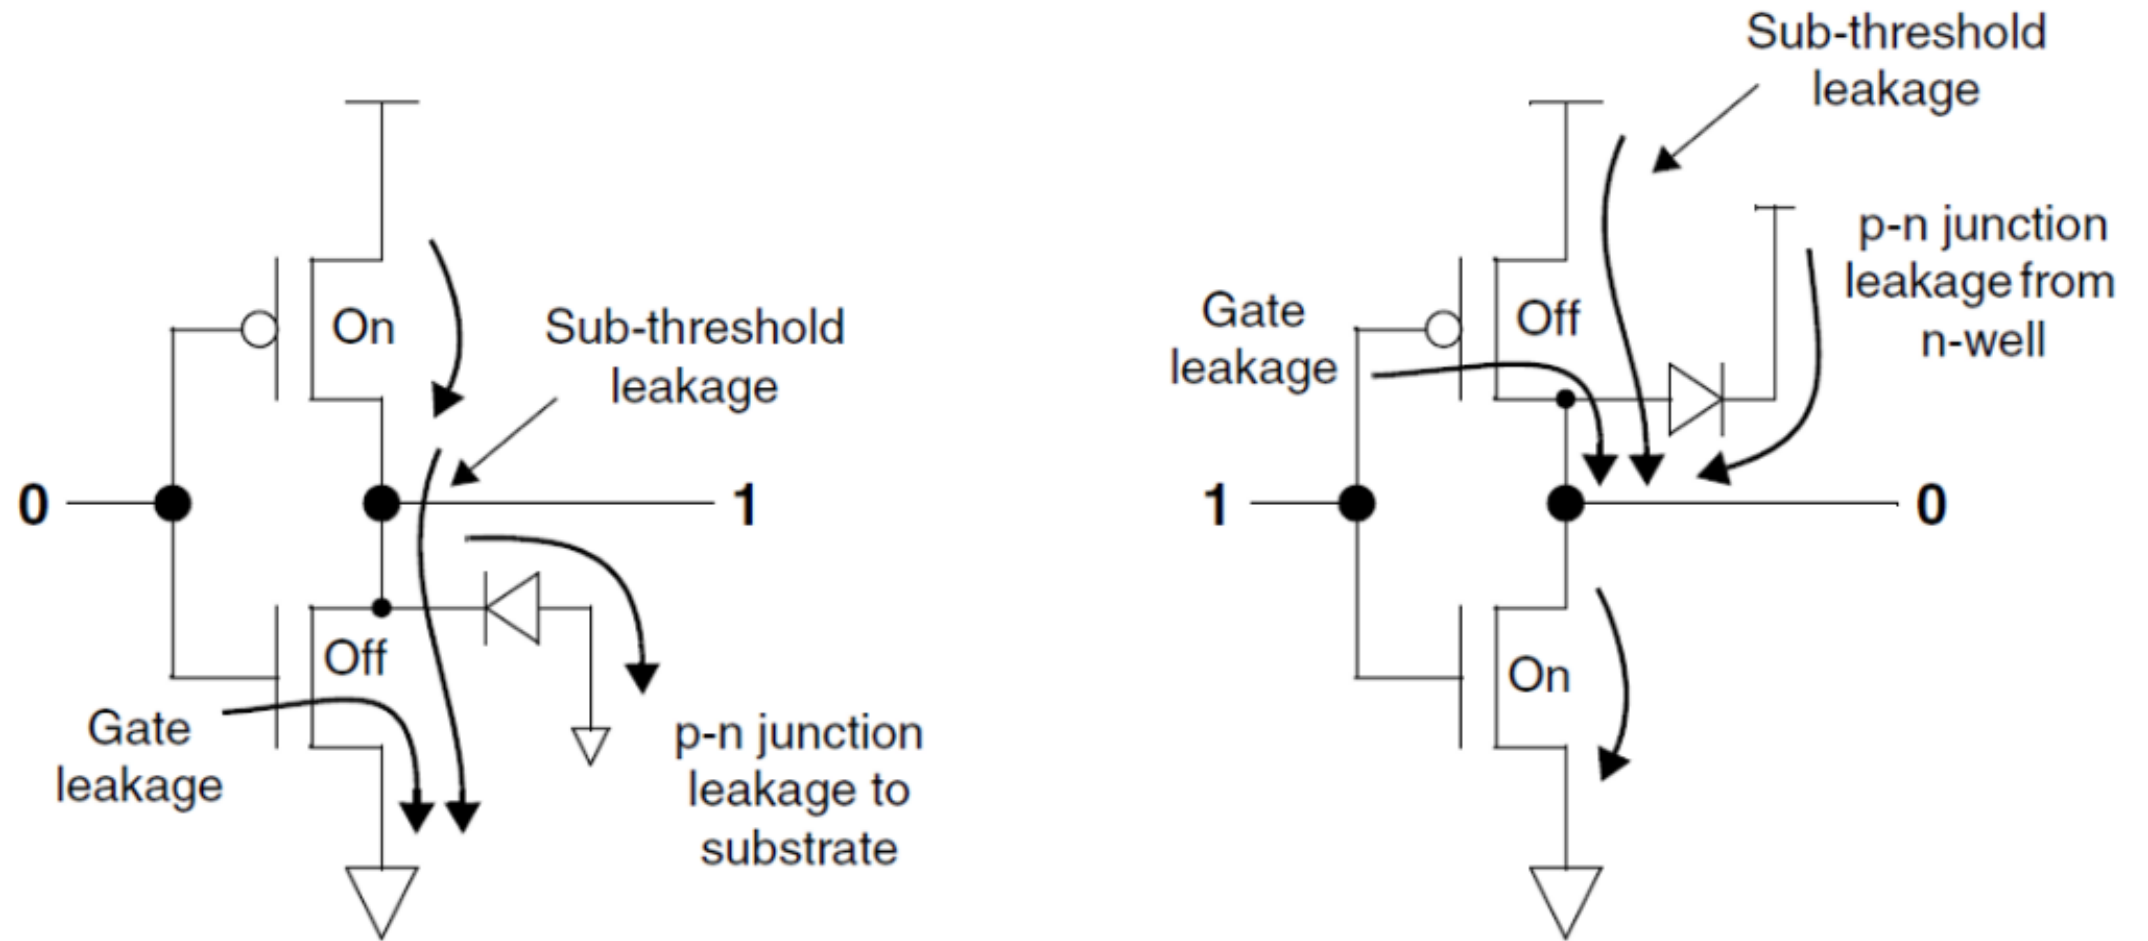


Figure 1.8: Static (Leakage) Power causes



RESULTS



Captured Data

- **Compute Metrics:** Metrics represent the counters that can be involved in a computation.
 - Number of additions.
 - Number of bytes transferred from the memory.
 - Number of instructions executed.
- **Power:** Core Voltage, Core Current, Memory Voltage, Memory Current.



Feature Selection

- Features represent the various groups of obtained data under the captured metrics.
- Reducing the feature count to minimize the complexity of the data, helps in efficient modelling.
- This can be achieved using largely available Machine Learning algorithms like PCA (or Principal Component Analysis) which helps eliminate the features from a set which do not contribute to large changes in the Target Variable (here, power) , i.e., feature prioritization.



Correlations in Data

- Multiple data points can be scaled versions of a data point.
- Such data points make data model complex and their removal will not affect the model in any way.
- Techniques to remove correlations:
 - Clustering.
 - Cosine Similarity.



Power Modelling

- Implemented Steps

- Minimum number of features obtained.
- Redundant data, representing similar data points removed.

1. Total Power from running workload: This is the power obtained from the GPU with operating workload.

2. Computed Power from Statistical Models: This is the computed power after Feature Selection and removing Correlations in data.

3. Accuracy of Prediction

- Implement Statistical Power model for energy consumed and obtain the convergence error.

- $Error = Total\ Energy - Computed\ Energy$



References

- [1] Siddhartha Sankar Mondal, “GPU: Power vs Performance”, Department of Information Technology, Uppsala Universitet, June 2014.
- [2] S Sundar, M Panchatcharam, “GPU Basics”, Department of Electronics and Communication, Indian Institute of Technology, August 2014, Madras.
- [3] Sunpyo Hong, Hyesoon Kim, “An Integrated GPU Power and Performance Model”, School of Computer Science, Georgia Institute of Technology, in International Parallel and Distributed Processing Symposium, USA, July 2010.
- [4] D.C Price, M.A Clark, “Optimizing performance-per-watt on GPUs in High Performance Computing”, *arXiv.org*, vol 3, 20th October 2015, pp 1-9.
- [5] Hitoshi Nagasaka, Naoya Maruyama, “Statistical Power Modelling of GPU Kernels using Performance Counters”, in *IEEE Explore, Japan Science and Technology Agency, CREST*, July 2010, pp 1-8.
- [6] CUDA – Introduction to GPU, www.tutorialspoint.com/cuda
- [7] GPU Fundamentals , <https://slideshare/GPU>



- [8] Y. Abe, H. Sasaki, S. Kato et al., “Power and performance characterization and modeling of gpu-accelerated systems,” in *Parallel and Distributed Processing Symposium, 2014 IEEE 28th International*, May 2014, pp. 113–122.
- [9] I. Paul, W. Huang, M. Arora et al., “Harmonia: Balancing compute and memory power in high-performance gpus,” in *Proceedings of the 42nd Annual International Symposium on Computer Architecture, ser. ISCA '15*. New York, NY, USA: ACM, 2015, pp. 54–65.
- [10] G. Wu, J. L. Greathouse, A. Lyashevsky et al., “Gpgpu performance and power estimation using machine learning,” in *2015 IEEE 21st International Symposium on High Performance Computer Architecture (HPCA)*, Feb 2015, pp. 564–576.
- [11] J. Coplin and M. Burtscher, “Energy, power, and performance characterization of gpgpu benchmark programs,” in *2016 IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW)*, May 2016, pp. 1190–1199.
- [12] G. Chen, B. Wu, D. Li, and X. Shen, “Purple: An extensible optimizer for portable data placement on gpu,” in *2014 47th Annual IEEE/ACM International Symposium on Microarchitecture*, Dec 2014, pp. 88–100.
- [13] R. Ge, R. Vogt, J. Majumder et al., “Effects of dynamic voltage and frequency scaling on a k20 gpu,” in *2013 42nd International Conference on Parallel Processing, Oct 2013*, pp. 826–833.



- [14] J. Guerreiro, A. Ilic, N. Roma et al., “Multi-kernel auto-tuning on gpus: Performance and energy-aware optimization,” in *2015 23rd Euromicro International Conference on Parallel, Distributed, and Network-Based Processing*, March 2015, pp. 438–445.
- [15] E. Lindholm, J. Nickolls, S. Oberman et al., “Nvidia tesla: A unified graphics and computing architecture,” in *IEEE Micro*, vol. 28, no. 2, pp. 39–55, Mar. 2012.

